



< Interview Prep Career GenAI Prompt Engg ChatGPT LLM >

Regularization in Deep Learning with Python Code



shubham.jain

Last Updated : 03 Dec, 2024



One of the most common challenges encountered by [data science](#) professionals is the issue of overfitting. Have you ever experienced a scenario where your [machine learning](#) model or regression model excelled on the training data but struggled to make accurate predictions on unseen test data? This is often a trade-off between high bias (underfitting) and high variance (overfitting), impacting the model's ability to generalize effectively. [Overfitting](#) can lead to poor performance on new data, especially in the presence of outliers or noise in the training set. If so, this article addresses your concerns by exploring techniques such as regularization in [deep learning](#) and essential concepts of bagging, boosting, and stacking in [ensemble learning](#) to improve model generalization.

New Feature



Make This Article *Fun and Interactive*

with Flash Cards and Quizzes!

Close

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

Show details

In this article, you will delve into the concept of regularization in [neural networks](#), discover different regularization techniques, and understand how these approaches can greatly enhance model performance while preventing overfitting.

Note: This article assumes that you have basic knowledge of neural networks and their implementation in Keras. If not, you can refer to the below articles first–

- [Fundamentals of Deep Learning – Starting with Artificial Neural Network](#)
- [Tutorial: Optimizing Neural Networks using Keras \(with Image recognition case study\)](#)

Table of contents

3. Different Regularization Techniques in Deep Learning

- L2&L1 Regularization
- Dropout
- Data Augmentation
- Early Stopping

4. A Case Study on MNIST Data with Keras

5. Conclusion

6. Frequently Asked Questions

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

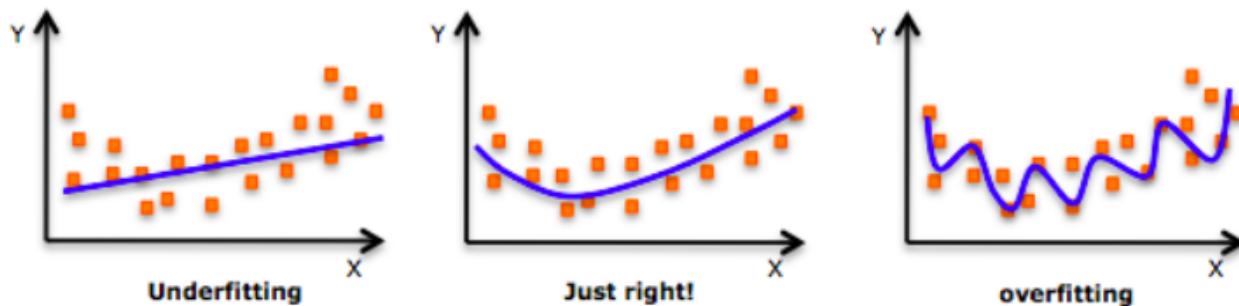
Show details

Regularization is a technique used in machine learning and deep learning to prevent overfitting and improve a model's generalization performance. It involves adding a penalty term to the [**loss function**](#) during training.

This penalty discourages the model from becoming too complex or having large parameter values, which helps in controlling the model's ability to fit noise in the training data. Regularization in deep learning [**methods**](#) includes L1 and L2 regularization, dropout, early stopping, and more. By applying regularization for deep learning, models become more robust and better at making accurate predictions on unseen data.

Also Read: [**Regularization in Machine Learning**](#)

Before we deep dive into the topic, take a look at this image:



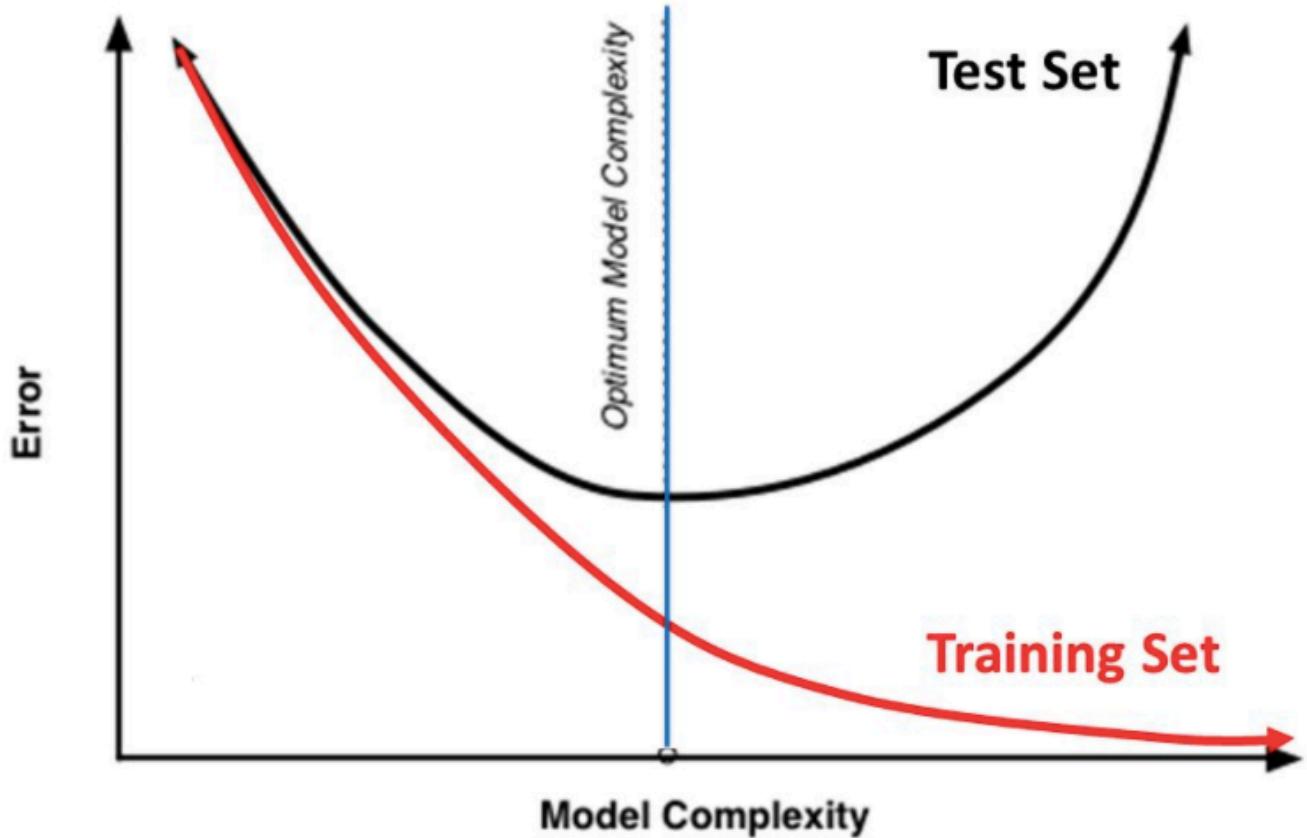
Have you seen this image before? As we move towards the right in this image, our model tries to learn too well the details and the noise from the training data, ultimately resulting in poor performance on the unseen data.

In other words, while going toward the right, the complexity of the model increases

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

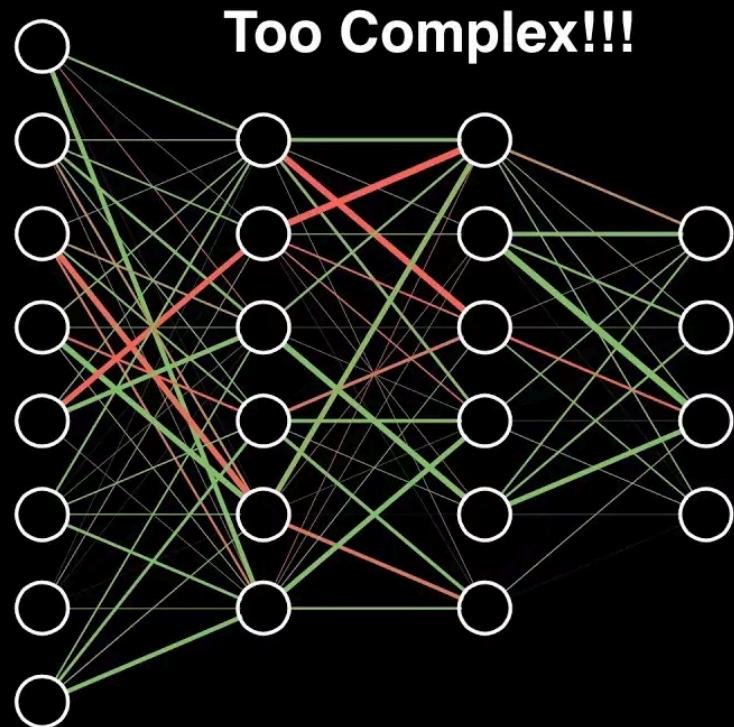
Training Vs. Test Set Error



If you've built a neural network before, you know how complex they are. This makes them more prone to overfitting.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details



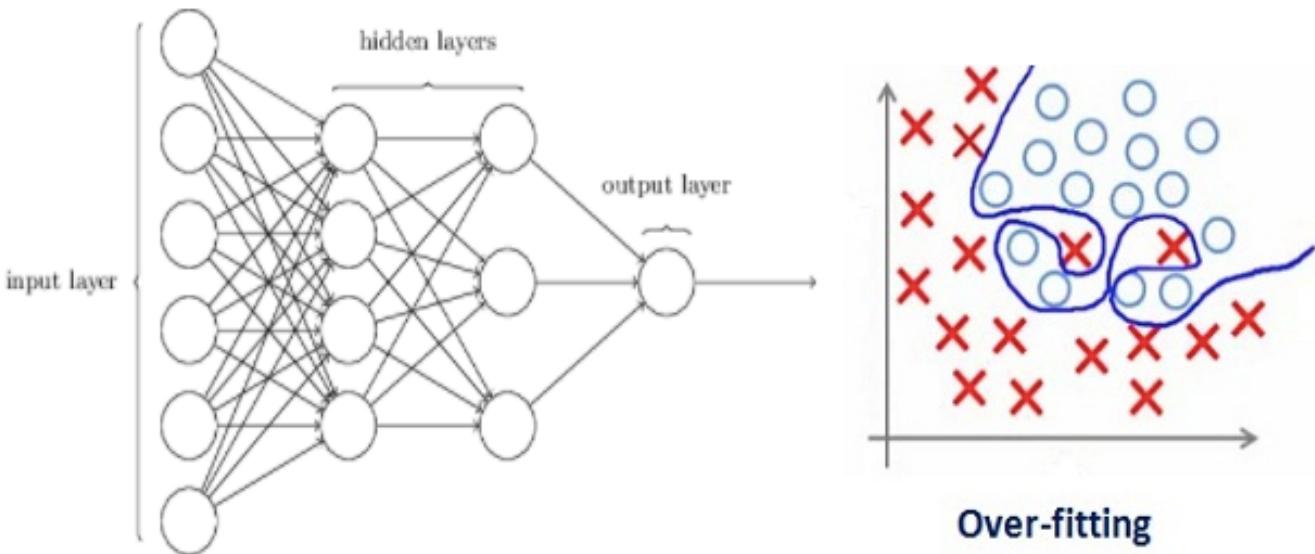
Regularization is a technique that modifies the learning algorithm slightly so that the model generalizes better. This, in turn, improves the model's performance on unseen data as well.

How does Regularization help Reduce Overfitting?

Let's consider a neural network that is overfitting on the training data as shown in the image below:

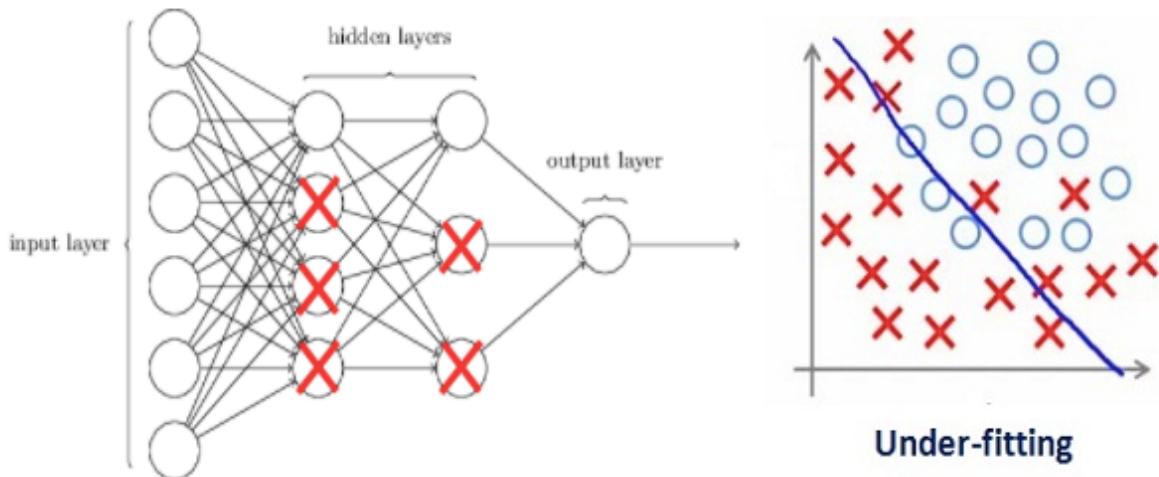
We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details



If you have studied the concept of regularization in machine learning, you will have a fair idea that regularization penalizes the coefficients. In deep learning, it penalizes the weight matrices of the nodes.

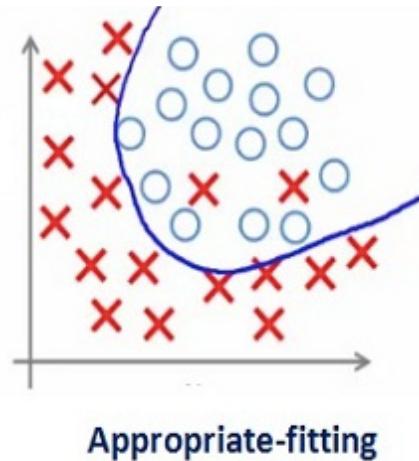
Assume that our regularization coefficient is so high that some of the weight matrices are nearly equal to zero.



We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

Such a large value of the regularization coefficient is not that useful. We need to optimize the value of the regularization coefficient to obtain a well-fitted model as shown in the image below:



Also Read: [How to avoid Over-fitting using Regularization?](#)

Different Regularization Techniques in Deep Learning

Now that we understand how regularization helps reduce overfitting, we'll learn a few different techniques for applying regularization in deep learning.

L2&L1 Regularization

L1 and L2 are the most common types of regularization in deep learning. These update the general cost function by adding another term known as the regularization term.

- *Cost function = Loss (say, binary cross entropy) + Regularization term*

Due to the addition of this regularization term, the values of weight matrices decrease

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

For L2:

$$\text{Cost function} = \text{Loss} + \frac{\lambda}{2m} * \sum \|w\|^2$$

For L1:

$$\text{Cost function} = \text{Loss} + \frac{\lambda}{2m} * \sum \|w\|$$

In L2, we have: $\|w\|^2 = \sum w_i^2$. This is known as ridge regression, where lambda is the regularization parameter. It is the hyperparameter whose value is optimized for better results. L2 regularization is also known as weight decay as it forces the weights to decay towards zero (but not exactly zero).

In L1, we having: $\|w\| = \sum |w_i|$. In this, we penalize the absolute value of the weights. Unlike L2, the weights may be reduced to zero here. L1 regularization is also called lasso regression. **Hence, it is very useful when we are trying to compress our model. Otherwise, we usually prefer L2 over it.**

In keras, we can directly apply regularization for deep learning to any layer using the [regularizers](#).

Below is the sample code to apply L2 regularization to a Dense layer:

```
from keras import regularizers

model.add(Dense(64, input_dim=64,
                kernel_regularizer=regularizers.l2(0.01)))
```

[Copy Code](#)

Note: Here the value 0.01 is the value of regularization parameter i.e. lambda which

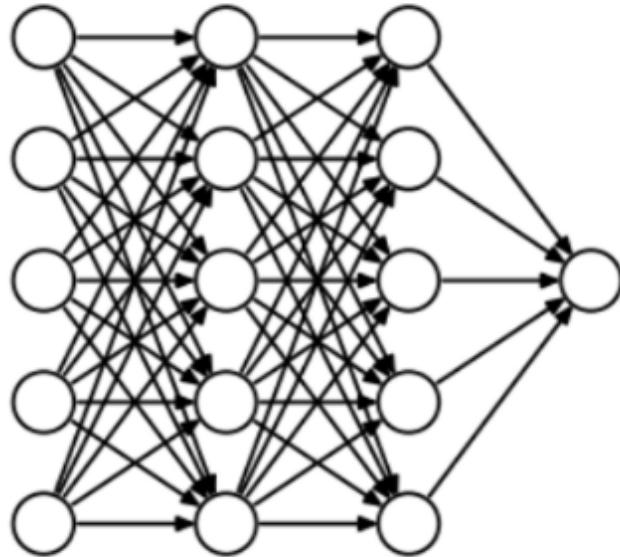
We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

Dropout

This is one of the most interesting types of [regularization techniques](#). It also produces very good results and is consequently the most frequently used regularization technique in the field of deep learning.

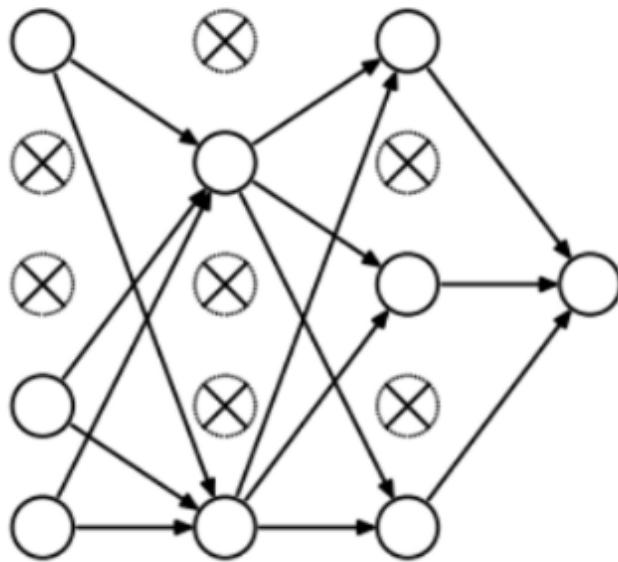
To understand dropout, let's say our neural network structure is akin to the one shown below:



So what does dropout do? At every iteration, it randomly selects some nodes and removes them along with all of their incoming and outgoing connections as shown below:

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

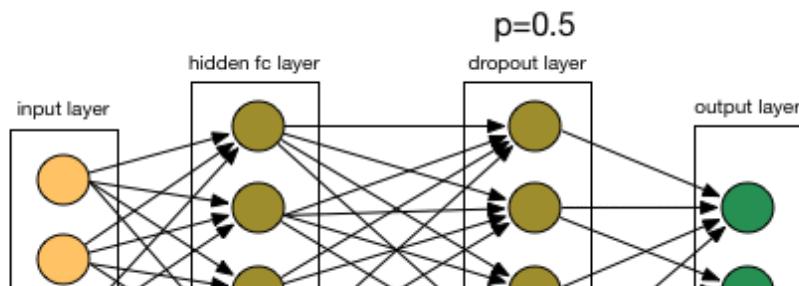
Show details



Each iteration has a different set of nodes, which results in a different set of outputs. This can also be thought of as an ensemble technique in machine learning.

Ensemble models usually perform better than a single model as they capture more randomness. Similarly, dropout models also perform better than normal neural network models.

This probability of choosing how many nodes should be dropped is the hyperparameter of the dropout function. As seen in the image above, dropout can be applied to both the hidden layers as well as the input layers.



We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

Due to these reasons, dropout is usually preferred when we have a large neural network structure to introduce more randomness.

In [Keras](#), we can implement dropout using the [Keras core layer](#). Below is the Python code for it:

```
from keras.layers.core import Dropout
```

Copy Code

```
model = Sequential([
    Dense(output_dim=hidden1_num_units, input_dim=input_num_units, activation='relu'),
    Dropout(0.25),

    Dense(output_dim=output_num_units, input_dim=hidden5_num_units, activation='softmax'),
])
```

As you can see, we have defined 0.25 as the probability of dropping. We can tune it further for better results using the grid search method.

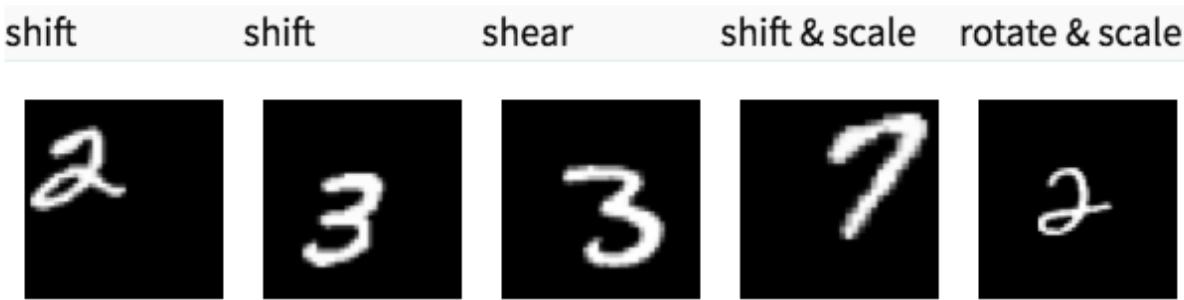
Data Augmentation

The simplest way to reduce overfitting is to increase the training data size. In machine learning, however, increasing the training data size was impossible as the labeled data was too costly.

But now, let's consider we are dealing with images. In this case, there are a few ways of increasing the size of the training data—rotating the image, flipping, scaling, shifting, etc. In the image below, some transformation has been done on the handwritten digits dataset.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

Show details



This technique is known as data augmentation. It usually provides a big leap in improving the accuracy of the model, and it can be considered a mandatory trick to improve our predictions.

In *keras*, we can perform all of these transformations using [ImageDataGenerator](#). It has a big list of arguments that you can use to pre-process your training data.

Below is the sample code to implement it:

```
from keras.preprocessing.image import ImageDataGenerator
datagen = ImageDataGenerator(horizontal_flip=True)
datagen.fit(train)
```

[Copy Code](#)

Early Stopping

Early stopping is a [cross-validation](#) strategy in which we keep one part of the training set as the validation set. When we see that the performance on the validation set is getting worse, we immediately stop the training on the model.



We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

In the above image, we will stop training at the dotted line since, after that, our model will start overfitting on the training data.

In *keras*, we can apply early stopping using the [**callbacks**](#) function. Below is the sample code for it.

```
from keras.callbacks import EarlyStopping  
  
EarlyStopping(monitor='val_err', patience=5)
```

[Copy Code](#)

Here, monitor refers to the quantity that you need to keep track of, and '**val_err**' refers to the validation error.

Patience denotes the number of epochs with no further improvement, after which training stops. For a better understanding, let's look at the above image again. After the dotted line, each epoch will result in a higher validation error value. Therefore, our model will stop 5 epochs after the dotted line (since our patience equals 5) because it sees no further improvement

Note: After 5 epochs (the value generally defined for patience), the model might start improving again, and the validation error may decrease. Therefore, we need to take extra care while tuning this hyperparameter.

Also Read: [**Prevent Overfitting Using Regularization Techniques**](#)

A Case Study on MNIST Data with Keras

By this point, you should theoretically understand the different techniques we have

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

[Copy Code](#)

```
%pylab inline

import numpy as np
import pandas as pd
from scipy.misc import imread
from sklearn.metrics import accuracy_score

from matplotlib import pyplot

import tensorflow as tf
import keras

# To stop potential randomness
seed = 128
rng = np.random.RandomState(seed)
```

Now, let's load the dataset.

```
root_dir = os.path.abspath('/Users/shubhamjain/Downloads/AV/identify the digits/')
data_dir = os.path.join(root_dir, 'data')
sub_dir = os.path.join(root_dir, 'sub')

## reading train file only
train = pd.read_csv(os.path.join(data_dir, 'Train', 'train.csv'))
train.head()
```

Output:

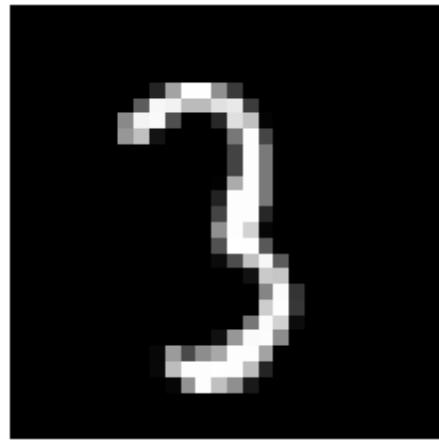
	filename	label
0	0.png	4
1	1.png	9
2	2.png	1
3	3.png	7
4	4.png	3

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

```
img = imread(filepath, flatten=True)

pylab.imshow(img, cmap='gray')
pylab.axis('off')
pylab.show()
```



```
#storing images in numpy arrays
temp = []
for img_name in train.filename:
    image_path = os.path.join(data_dir, 'Train', 'Images', 'train', img_name)
    img = imread(image_path, flatten=True)
    img = img.astype('float32')
    temp.append(img)

x_train = np.stack(temp)

x_train /= 255.0
x_train = x_train.reshape(-1, 784).astype('float32')

y_train = keras.utils.np_utils.to_categorical(train.label.values)
```

[Copy Code](#)

Create a validation dataset to optimize our model for better scores. We will use a 70:30 train-validation dataset ratio.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

First, let's build a simple neural network with 5 hidden layers, each with 500 nodes.

```
# import keras modules
from keras.models import Sequential
from keras.layers import Dense

# define vars
input_num_units = 784
hidden1_num_units = 500
hidden2_num_units = 500
hidden3_num_units = 500
hidden4_num_units = 500
hidden5_num_units = 500
output_num_units = 10

epochs = 10
batch_size = 128

model = Sequential([
    Dense(output_dim=hidden1_num_units, input_dim=input_num_units, activation='relu'),
    Dense(output_dim=hidden2_num_units, input_dim=hidden1_num_units, activation='relu'),
    Dense(output_dim=hidden3_num_units, input_dim=hidden2_num_units, activation='relu'),
    Dense(output_dim=hidden4_num_units, input_dim=hidden3_num_units, activation='relu'),
    Dense(output_dim=hidden5_num_units, input_dim=hidden4_num_units, activation='relu'),
    Dense(output_dim=output_num_units, input_dim=hidden5_num_units, activation='softmax'),
])
```

[Copy Code](#)

Also Read: [Train-Test-Validation Split in 2024](#)

Evaluating the Model

Note that we are just running it for 10 epochs. Let's quickly check the performance of our model.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

Output:

```
Train on 34300 samples, validate on 14700 samples
Epoch 1/10
34300/34300 [=====] - 11s - loss: 0.2926 - acc: 0.9118 - val_loss: 0.1854 - val_acc: 0.9444
Epoch 2/10
34300/34300 [=====] - 10s - loss: 0.1145 - acc: 0.9643 - val_loss: 0.1279 - val_acc: 0.9627
Epoch 3/10
34300/34300 [=====] - 10s - loss: 0.0792 - acc: 0.9765 - val_loss: 0.1114 - val_acc: 0.9671
Epoch 4/10
34300/34300 [=====] - 10s - loss: 0.0616 - acc: 0.9814 - val_loss: 0.1213 - val_acc: 0.9671
Epoch 5/10
34300/34300 [=====] - 10s - loss: 0.0468 - acc: 0.9855 - val_loss: 0.1280 - val_acc: 0.9665
Epoch 6/10
34300/34300 [=====] - 10s - loss: 0.0379 - acc: 0.9882 - val_loss: 0.1299 - val_acc: 0.9665
Epoch 7/10
34300/34300 [=====] - 10s - loss: 0.0404 - acc: 0.9874 - val_loss: 0.1203 - val_acc: 0.9692
Epoch 8/10
34300/34300 [=====] - 11s - loss: 0.0351 - acc: 0.9897 - val_loss: 0.1200 - val_acc: 0.9716
Epoch 9/10
34300/34300 [=====] - 10s - loss: 0.0263 - acc: 0.9918 - val_loss: 0.1089 - val_acc: 0.9736
Epoch 10/10
34300/34300 [=====] - 12s - loss: 0.0237 - acc: 0.9931 - val_loss: 0.1308 - val_acc: 0.9699
```

Now, let's try the L2 regularizer over it and check whether it gives better results than a simple neural network model.

```
from keras import regularizers
```

[Copy Code](#)

```
model = Sequential([
    Dense(output_dim=hidden1_num_units, input_dim=input_num_units, activation='relu',
           kernel_regularizer=regularizers.l2(0.0001)),
    Dense(output_dim=hidden2_num_units, input_dim=hidden1_num_units, activation='relu',
           kernel_regularizer=regularizers.l2(0.0001)),
    Dense(output_dim=hidden3_num_units, input_dim=hidden2_num_units, activation='relu',
           kernel_regularizer=regularizers.l2(0.0001)),
    Dense(output_dim=hidden4_num_units, input_dim=hidden3_num_units, activation='relu',
           kernel_regularizer=regularizers.l2(0.0001)),
    Dense(output_dim=hidden5_num_units, input_dim=hidden4_num_units, activation='relu',
           kernel_regularizer=regularizers.l2(0.0001)),

    Dense(output_dim=output_num_units, input_dim=hidden5_num_units, activation='softmax'),
])

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

```
Train on 34300 samples, validate on 14700 samples
Epoch 1/10
34300/34300 [=====] - 15s - loss: 0.4862 - acc: 0.9092 - val_loss: 0.3283 - val_acc: 0.9531
Epoch 2/10
34300/34300 [=====] - 14s - loss: 0.2887 - acc: 0.9635 - val_loss: 0.2822 - val_acc: 0.9644
Epoch 3/10
34300/34300 [=====] - 14s - loss: 0.2388 - acc: 0.9738 - val_loss: 0.2534 - val_acc: 0.9696
Epoch 4/10
34300/34300 [=====] - 12s - loss: 0.2010 - acc: 0.9806 - val_loss: 0.2418 - val_acc: 0.9678
Epoch 5/10
34300/34300 [=====] - 13s - loss: 0.1807 - acc: 0.9830 - val_loss: 0.2470 - val_acc: 0.9655
Epoch 6/10
34300/34300 [=====] - 13s - loss: 0.1707 - acc: 0.9826 - val_loss: 0.2458 - val_acc: 0.9669
Epoch 7/10
34300/34300 [=====] - 12s - loss: 0.1524 - acc: 0.9865 - val_loss: 0.2125 - val_acc: 0.9718
Epoch 8/10
34300/34300 [=====] - 14s - loss: 0.1365 - acc: 0.9886 - val_loss: 0.2254 - val_acc: 0.9697
Epoch 9/10
34300/34300 [=====] - 13s - loss: 0.1392 - acc: 0.9863 - val_loss: 0.2070 - val_acc: 0.9695
Epoch 10/10
34300/34300 [=====] - 13s - loss: 0.1209 - acc: 0.9912 - val_loss: 0.1952 - val_acc: 0.9739
```

Note that the value of **lambda** is equal to 0.0001. Great! We just obtained an accuracy that is greater than our previous NN model.

Now, let's try the L1 regularization technique.

l1

[Copy Code](#)

```
model = Sequential([
    Dense(output_dim=hidden1_num_units, input_dim=input_num_units, activation='relu',
           kernel_regularizer=regularizers.l1(0.0001)),
    Dense(output_dim=hidden2_num_units, input_dim=hidden1_num_units, activation='relu',
           kernel_regularizer=regularizers.l1(0.0001)),
    Dense(output_dim=hidden3_num_units, input_dim=hidden2_num_units, activation='relu',
           kernel_regularizer=regularizers.l1(0.0001)),
    Dense(output_dim=hidden4_num_units, input_dim=hidden3_num_units, activation='relu',
           kernel_regularizer=regularizers.l1(0.0001)),
    Dense(output_dim=hidden5_num_units, input_dim=hidden4_num_units, activation='relu',
           kernel_regularizer=regularizers.l1(0.0001)),

    Dense(output_dim=output_num_units, input_dim=hidden5_num_units, activation='softmax'),
])

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
trained_model_5d = model.fit(x_train, y_train, nb_epoch=epochs, batch_size=batch_size, v:
```

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

```
Train on 34300 samples, validate on 14700 samples
Epoch 1/10
34300/34300 [=====] - 13s - loss: 2.8190 - acc: 0.9019 - val_loss: 1.7565 - val_acc: 0.9455
Epoch 2/10
34300/34300 [=====] - 13s - loss: 1.3347 - acc: 0.9541 - val_loss: 1.0242 - val_acc: 0.9565
Epoch 3/10
34300/34300 [=====] - 12s - loss: 0.8235 - acc: 0.9628 - val_loss: 0.7109 - val_acc: 0.9554
Epoch 4/10
34300/34300 [=====] - 15s - loss: 0.5809 - acc: 0.9661 - val_loss: 0.5493 - val_acc: 0.9580
Epoch 5/10
34300/34300 [=====] - 15s - loss: 0.4472 - acc: 0.9721 - val_loss: 0.4364 - val_acc: 0.9637
Epoch 6/10
34300/34300 [=====] - 14s - loss: 0.3720 - acc: 0.9754 - val_loss: 0.3725 - val_acc: 0.9659
Epoch 7/10
34300/34300 [=====] - 16s - loss: 0.3208 - acc: 0.9778 - val_loss: 0.3323 - val_acc: 0.9690
Epoch 8/10
34300/34300 [=====] - 14s - loss: 0.2893 - acc: 0.9790 - val_loss: 0.3114 - val_acc: 0.9688
Epoch 9/10
34300/34300 [=====] - 16s - loss: 0.2619 - acc: 0.9820 - val_loss: 0.3068 - val_acc: 0.9673
Epoch 10/10
34300/34300 [=====] - 15s - loss: 0.2418 - acc: 0.9828 - val_loss: 0.2921 - val_acc: 0.9667
```

This doesn't show any improvement over the previous model. Let's jump to the dropout technique.

```
## dropout
```

[Copy Code](#)

```
from keras.layers.core import Dropout
model = Sequential([
    Dense(output_dim=hidden1_num_units, input_dim=input_num_units, activation='relu'),
    Dropout(0.25),
    Dense(output_dim=hidden2_num_units, input_dim=hidden1_num_units, activation='relu'),
    Dropout(0.25),
    Dense(output_dim=hidden3_num_units, input_dim=hidden2_num_units, activation='relu'),
    Dropout(0.25),
    Dense(output_dim=hidden4_num_units, input_dim=hidden3_num_units, activation='relu'),
    Dropout(0.25),
    Dense(output_dim=hidden5_num_units, input_dim=hidden4_num_units, activation='relu'),
    Dropout(0.25),

    Dense(output_dim=output_num_units, input_dim=hidden5_num_units, activation='softmax'),
])

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
trained_model_5d = model.fit(x_train, y_train, nb_epoch=epochs, batch_size=batch_size, v:
```

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

```

Dropout(0.25),
Dense(output_dim=hidden2_num_units, input_dim=hidden1_num_units, activation='relu'),
Dropout(0.25),
Dense(output_dim=hidden3_num_units, input_dim=hidden2_num_units, activation='relu'),
Dropout(0.25),
Dense(output_dim=hidden4_num_units, input_dim=hidden3_num_units, activation='relu'),
Dropout(0.25),
Dense(output_dim=hidden5_num_units, input_dim=hidden4_num_units, activation='relu'),
Dropout(0.25),

Dense(output_dim=output_num_units, input_dim=hidden5_num_units, activation='softmax'),
])

```

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['acc'])
 trained_model_5d = model.fit(x_train, y_train, nb_epoch=epochs, batch_size=batch_size, v:

Output:

```

Train on 34300 samples, validate on 14700 samples
Epoch 1/10
34300/34300 [=====] - 14s - loss: 0.4231 - acc: 0.8652 - val_loss: 0.1756 - val_acc: 0.9461
Epoch 2/10
34300/34300 [=====] - 12s - loss: 0.1771 - acc: 0.9487 - val_loss: 0.1526 - val_acc: 0.9573
Epoch 3/10
34300/34300 [=====] - 13s - loss: 0.1296 - acc: 0.9625 - val_loss: 0.1111 - val_acc: 0.9683
Epoch 4/10
34300/34300 [=====] - 13s - loss: 0.1055 - acc: 0.9691 - val_loss: 0.1137 - val_acc: 0.9688
Epoch 5/10
34300/34300 [=====] - 12s - loss: 0.0924 - acc: 0.9739 - val_loss: 0.1184 - val_acc: 0.9677
Epoch 6/10
34300/34300 [=====] - 12s - loss: 0.0821 - acc: 0.9755 - val_loss: 0.1059 - val_acc: 0.9712
Epoch 7/10
34300/34300 [=====] - 13s - loss: 0.0714 - acc: 0.9794 - val_loss: 0.1103 - val_acc: 0.9699
Epoch 8/10
34300/34300 [=====] - 13s - loss: 0.0635 - acc: 0.9812 - val_loss: 0.0965 - val_acc: 0.9766
Epoch 9/10
34300/34300 [=====] - 13s - loss: 0.0588 - acc: 0.9825 - val_loss: 0.1133 - val_acc: 0.9718
Epoch 10/10
34300/34300 [=====] - 15s - loss: 0.0572 - acc: 0.9826 - val_loss: 0.1085 - val_acc: 0.9729

```

Not bad! Dropout also gives us a little improvement over our simple NN model.

Now, let's try data augmentation.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

```

temp = []
for img_name in train.filename:
    image_path = os.path.join(data_dir, 'Train', 'Images', 'train', img_name)
    img = imread(image_path, flatten=True)
    img = img.astype('float32')
    temp.append(img)

x_train = np.stack(temp)

X_train = x_train.reshape(x_train.shape[0], 1, 28, 28)

X_train = X_train.astype('float32')

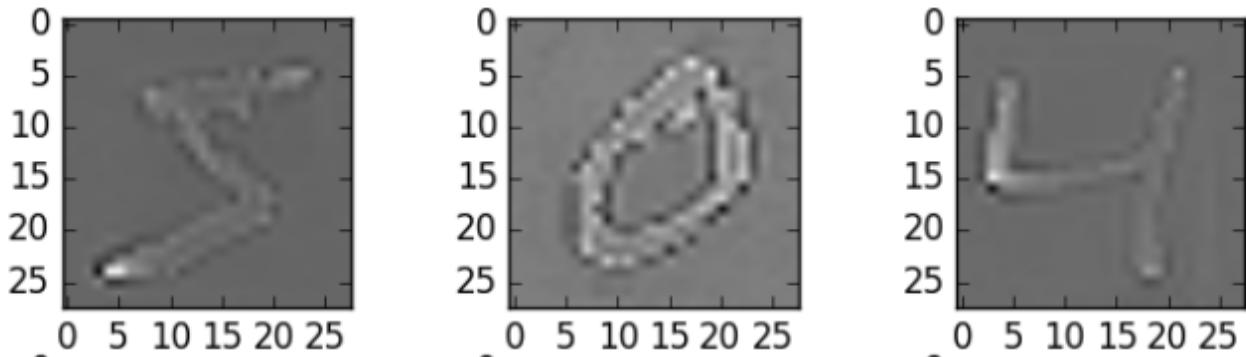
```

Now, fit the training data in order to augment.

```
# fit parameters from data
datagen.fit(X_train)
```

[Copy Code](#)

Here, I have used **zca_whitening** as the argument, which highlights the outline of each digit as shown in the image below:



```

## splitting
y_train = keras.utils.to_categorical(train.label.values)
split_size = int(x_train.shape[0]*0.7)

x_train, x_test = X_train[:split_size], X_train[split_size:]

```

[Copy Code](#)

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

[Copy Code](#)

```
## structure using dropout
from keras.layers.core import Dropout
model = Sequential([
    Dense(output_dim=hidden1_num_units, input_dim=input_num_units, activation='relu'),
    Dropout(0.25),
    Dense(output_dim=hidden2_num_units, input_dim=hidden1_num_units, activation='relu'),
    Dropout(0.25),
    Dense(output_dim=hidden3_num_units, input_dim=hidden2_num_units, activation='relu'),
    Dropout(0.25),
    Dense(output_dim=hidden4_num_units, input_dim=hidden3_num_units, activation='relu'),
    Dropout(0.25),
    Dense(output_dim=hidden5_num_units, input_dim=hidden4_num_units, activation='relu'),
    Dropout(0.25),

    Dense(output_dim=output_num_units, input_dim=hidden5_num_units, activation='softmax'),
])
```

```
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['acc'])  
trained_model_5d = model.fit(x_train, y_train, nb_epoch=epochs, batch_size=batch_size, v  
isuals=False)
```

Output:

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 20s - loss: 0.1222 - acc: 0.9669 - val_loss: 0.0893 - val_acc: 0.9761
Epoch 2/10
60000/60000 [=====] - 20s - loss: 0.0949 - acc: 0.9741 - val_loss: 0.0948 - val_acc: 0.9746
Epoch 3/10
60000/60000 [=====] - 19s - loss: 0.0849 - acc: 0.9764 - val_loss: 0.0793 - val_acc: 0.9780
Epoch 4/10
60000/60000 [=====] - 19s - loss: 0.0705 - acc: 0.9803 - val_loss: 0.0825 - val_acc: 0.9794
Epoch 5/10
60000/60000 [=====] - 19s - loss: 0.0671 - acc: 0.9809 - val_loss: 0.0765 - val_acc: 0.9782
Epoch 6/10
60000/60000 [=====] - 20s - loss: 0.0620 - acc: 0.9831 - val_loss: 0.0797 - val_acc: 0.9817
Epoch 7/10
60000/60000 [=====] - 18s - loss: 0.0563 - acc: 0.9843 - val_loss: 0.0746 - val_acc: 0.9828
Epoch 8/10
60000/60000 [=====] - 18s - loss: 0.0492 - acc: 0.9860 - val_loss: 0.0847 - val_acc: 0.9800
Epoch 9/10
60000/60000 [=====] - 18s - loss: 0.0515 - acc: 0.9855 - val_loss: 0.0880 - val_acc: 0.9784
Epoch 10/10
60000/60000 [=====] - 20s - loss: 0.0503 - acc: 0.9856 - val_loss: 0.0614 - val_acc: 0.9835
```

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)

```
from keras.callbacks import EarlyStopping
```

[Copy Code](#)

```
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=[ 'accuracy'])
trained_model_5d = model.fit(x_train, y_train, nb_epoch=epochs, batch_size=batch_size, verbose=1,
 , callbacks = [EarlyStopping(monitor='val_acc', patience=2)])
```

Output:

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 20s - loss: 0.0479 - acc: 0.9869 - val_loss: 0.0881 - val_acc: 0.9808
Epoch 2/10
60000/60000 [=====] - 19s - loss: 0.0486 - acc: 0.9875 - val_loss: 0.0756 - val_acc: 0.9833
Epoch 3/10
60000/60000 [=====] - 21s - loss: 0.0467 - acc: 0.9875 - val_loss: 0.0788 - val_acc: 0.9795
Epoch 4/10
60000/60000 [=====] - 20s - loss: 0.0398 - acc: 0.9886 - val_loss: 0.0821 - val_acc: 0.9818
Epoch 5/10
60000/60000 [=====] - 20s - loss: 0.0417 - acc: 0.9887 - val_loss: 0.0968 - val_acc: 0.9813
```

Our model stops after only 5 iterations as the validation accuracy does not improve. It gives good results in cases where we run it for a larger number of epochs, so it's a technique to optimize the number of epochs.

Also Read: [Techniques To Prevent Overfitting In Neural Networks](#)

Conclusion

I hope you now understand regularization deep learning techniques in deep learning and the different methods required to implement them, not just for linear regression models but also for complex neural networks like CNN. Regularization is crucial when training deep learning models on large datasets to prevent overfitting and improve generalization. You can effectively employ techniques like L1/L2 regularization, dropout, data augmentation, and early stopping using gradient descent optimization.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

learning methods when dealing with a deep learning task, as they will help you expand your horizons, better understand the topic, and build more robust machine learning algorithms.

Hope you like the article and better understand regularization in deep learning.

Regularisation techniques, such as L1 and L2 regularization, dropout, and data augmentation, prevent overfitting, and improve model performance.



[shubham.jain](#)

Deep Learning

Image

Intermediate

Python

Python

Technique

Unstructured Data

Free Courses



Generative AI - A Way of Life

Explore Generative AI for beginners: create text and images, use top AI tools, learn practical skills, and ethics.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details



Getting Started with Large Language Models

Master Large Language Models (LLMs) with this course, offering clear guidance in NLP and model training made simple.



Building LLM Applications using Prompt Engineering

This free course guides you on building LLM apps, mastering prompt engineering, and developing chatbots with enterprise data.



Improving Real World RAG Systems: Key Challenges & Practical Solutions

Explore practical solutions, advanced retrieval strategies, and agentic RAG systems to improve context, relevance, and accuracy in AI-driven applications.



We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

Master MS Excel for data analysis with key formulas, functions, and LookUp tools in this comprehensive course.

RECOMMENDED ARTICLES

[Ridge and Lasso Regression in Python](#)

[Regularization in Machine Learning](#)

[Dropout Regularization in Deep Learning](#)

[Complete Guide to Prevent Overfitting in Neural...](#)

[Complete Guide to Prevent Overfitting in Neural...](#)

[Understanding Overfitting in ConvNets](#)

[Prevent Overfitting Using Regularization Techni...](#)

[Tutorial: Optimizing Neural Networks using Kera...](#)

[Study of Regularization Techniques of Linear Mo...](#)

Responses From Readers

What are your thoughts?...

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

[Submit reply](#)

Pramod

Thanks a lot bro. Immensely helpful.



Chetan

Could we achieve the same validation accuracy with much lesser architecture ? If so, then our Hyperparameters (especially the L2 lambda) must be quite high (say 0.05). Isn't it ? Or, is there any other way of combating this ??



Florida Meacci

Thank you for the great article. I'm trying to differentiate

Frequently Asked Questions

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)

such as L1 and L2 regularization, dropout, and batch normalization help control model complexity and improve neural network generalization to unseen data.

Q2. What is L1 regularization and L2 regularization?

Q3. What is dropout in neural network?

Q4.What is L1 and L2 regularization technique?

Write for us →

Write, captivate, and earn accolades and rewards for your work

- Reach a Global Audience
- Get Expert Feedback
- Build Your Brand & Audience
- Cash In on Your Knowledge
- Join a Thriving Community
- Level Up Your Data Science Game

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details



Flagship Courses

GenAI Pinnacle Program | AI/ML BlackBelt Courses

Free Courses

Generative AI | Large Language Models | Building LLM Applications using Prompt Engineering | Building Your first RAG System using LlamaIndex | Stability.AI | MidJourney | Building Production Ready RAG systems using LlamaIndex | Building LLMs for Code | Deep Learning | Python | Microsoft Excel | Machine Learning | Decision Trees | Pandas for Data Analysis | Ensemble Learning | NLP | NLP using Deep Learning | Neural Networks | Loan Prediction Practice Problem | Time Series Forecasting | Tableau | Business Analytics

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

Generative AI Tools and Techniques

GANs | VAEs | Transformers | StyleGAN | Pix2Pix | Autoencoders | GPT | BERT | Word2Vec | LSTM | Attention Mechanisms | Diffusion Models | LLMs | SLMs | StyleGAN | Encoder Decoder Models | Prompt Engineering | LangChain | LlamaIndex | RAG | Fine-tuning | LangChain AI Agent | Multimodal Models | RNNs | DCGAN | ProGAN | Text-to-Image Models | DDPM | Document Question Answering | Imagen | T5 (Text-to-Text Transfer Transformer) | Seq2seq Models | WaveNet | Attention Is All You Need (Transformer Architecture)

Popular GenAI Models

Llama 3.1 | Llama 3 | Llama 2 | GPT 4o Mini | GPT 4o | GPT 3 | Claude 3 Haiku | Claude 3.5 Sonnet | Phi 3.5 | Phi 3 | Mistral Large 2 | Mistral NeMo | Mistral-7b | Gemini 1.5 Pro | Gemini Flash 1.5 | Bedrock | Vertex AI | DALL.E | Midjourney | Stable Diffusion

Data Science Tools and Techniques

Python | R | SQL | Jupyter Notebooks | TensorFlow | Scikit-learn | PyTorch | Tableau | Apache Spark | Matplotlib | Seaborn | Pandas | Hadoop | Docker | Git | Keras | Apache Kafka | AWS | NLP | Random Forest | Computer Vision | Data Visualization | Data Exploration | Big Data | Common Machine Learning Algorithms | Machine Learning

Company

About Us

Contact Us

Careers

Discover

Blogs

Expert session

Podcasts

Comprehensive Guides

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

[GenAI Program](#)[Agentic AI Pioneer Program](#)**Contribute**[Become an Author](#)[Become a speaker](#)[Become a mentor](#)[Become an instructor](#)[Events](#)[AI Newsletter](#)**Enterprise**[Our offerings](#)[Trainings](#)[Data Culture](#)

[Terms & conditions](#) • [Refund Policy](#) • [Privacy Policy](#) • [Cookies Policy](#) © Analytics Vidhya
2025. All rights reserved.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)